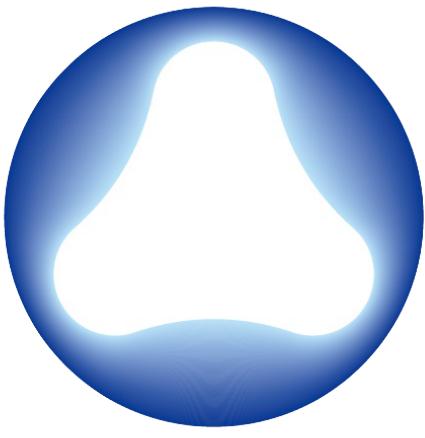




ICLR



RaSA: Rank-Sharing Low-Rank Adaptation

Zhiwei He Zhaopeng Tu Xing Wang Xingyu Chen Zhijie Wang
Jiahao Xu Tian Liang Wenxiang Jiao Zhuosheng Zhang Rui Wang

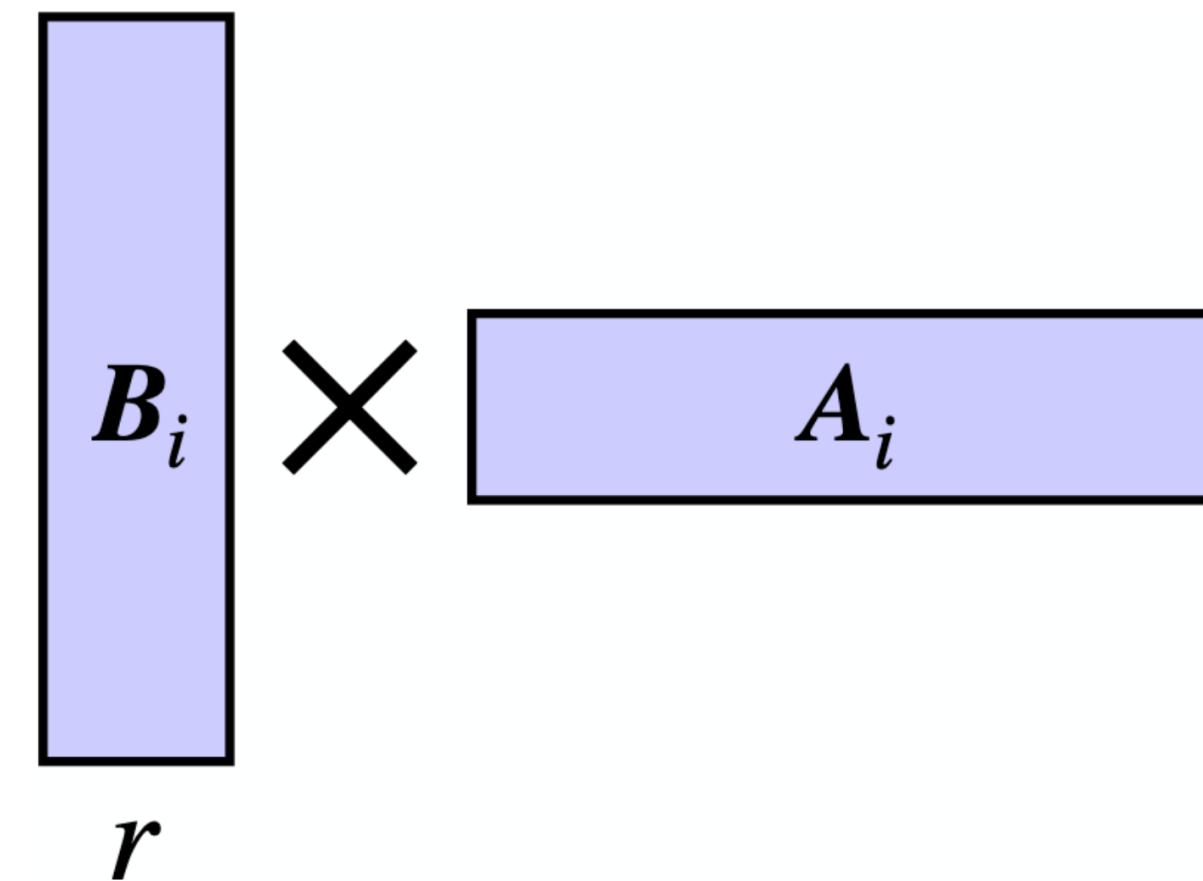
Presenter: Zhiwei He

Low-Rank Adaptation (LoRA)

Low-rank update

- Parameter Update:

$$\mathbf{W}_i + \underbrace{\Delta \mathbf{W}_i}_{\text{low-rank}} =$$

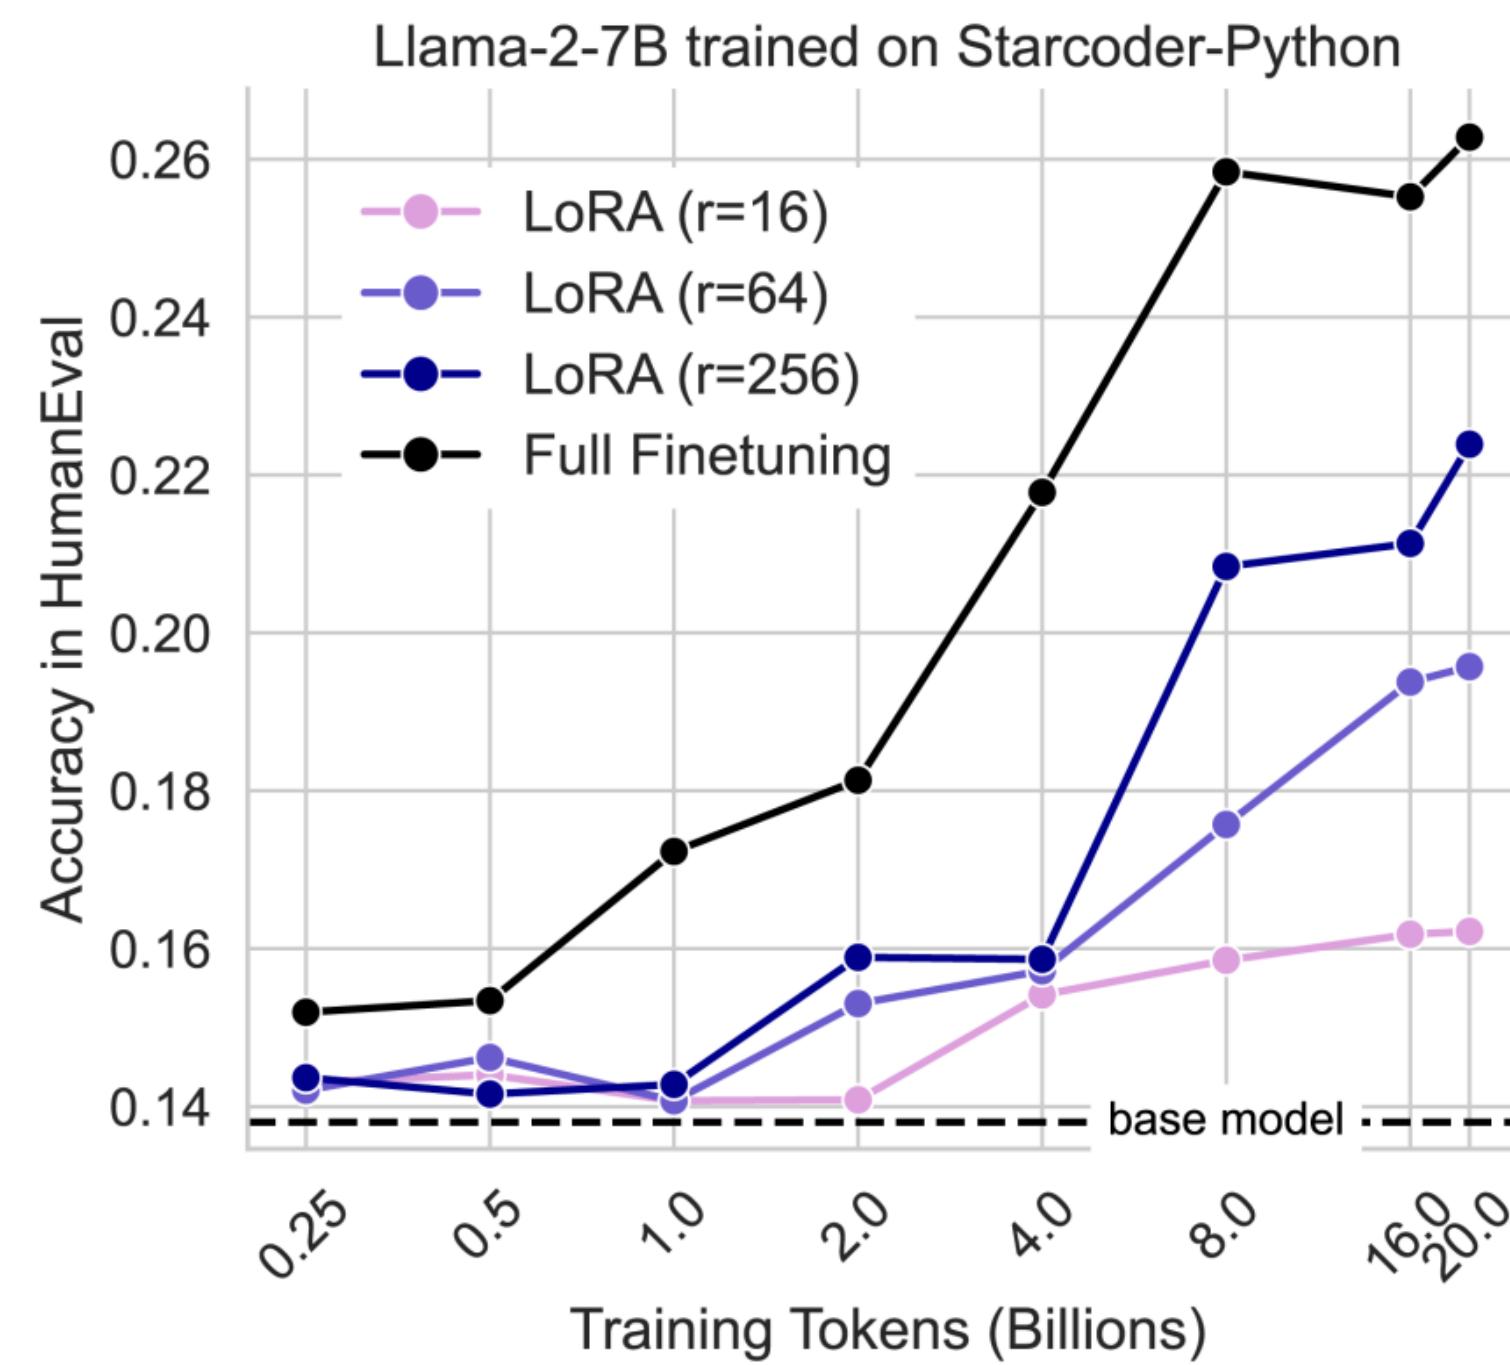


$$B_i \in \mathbb{R}^{b \times r}, A_i \in \mathbb{R}^{r \times a}$$

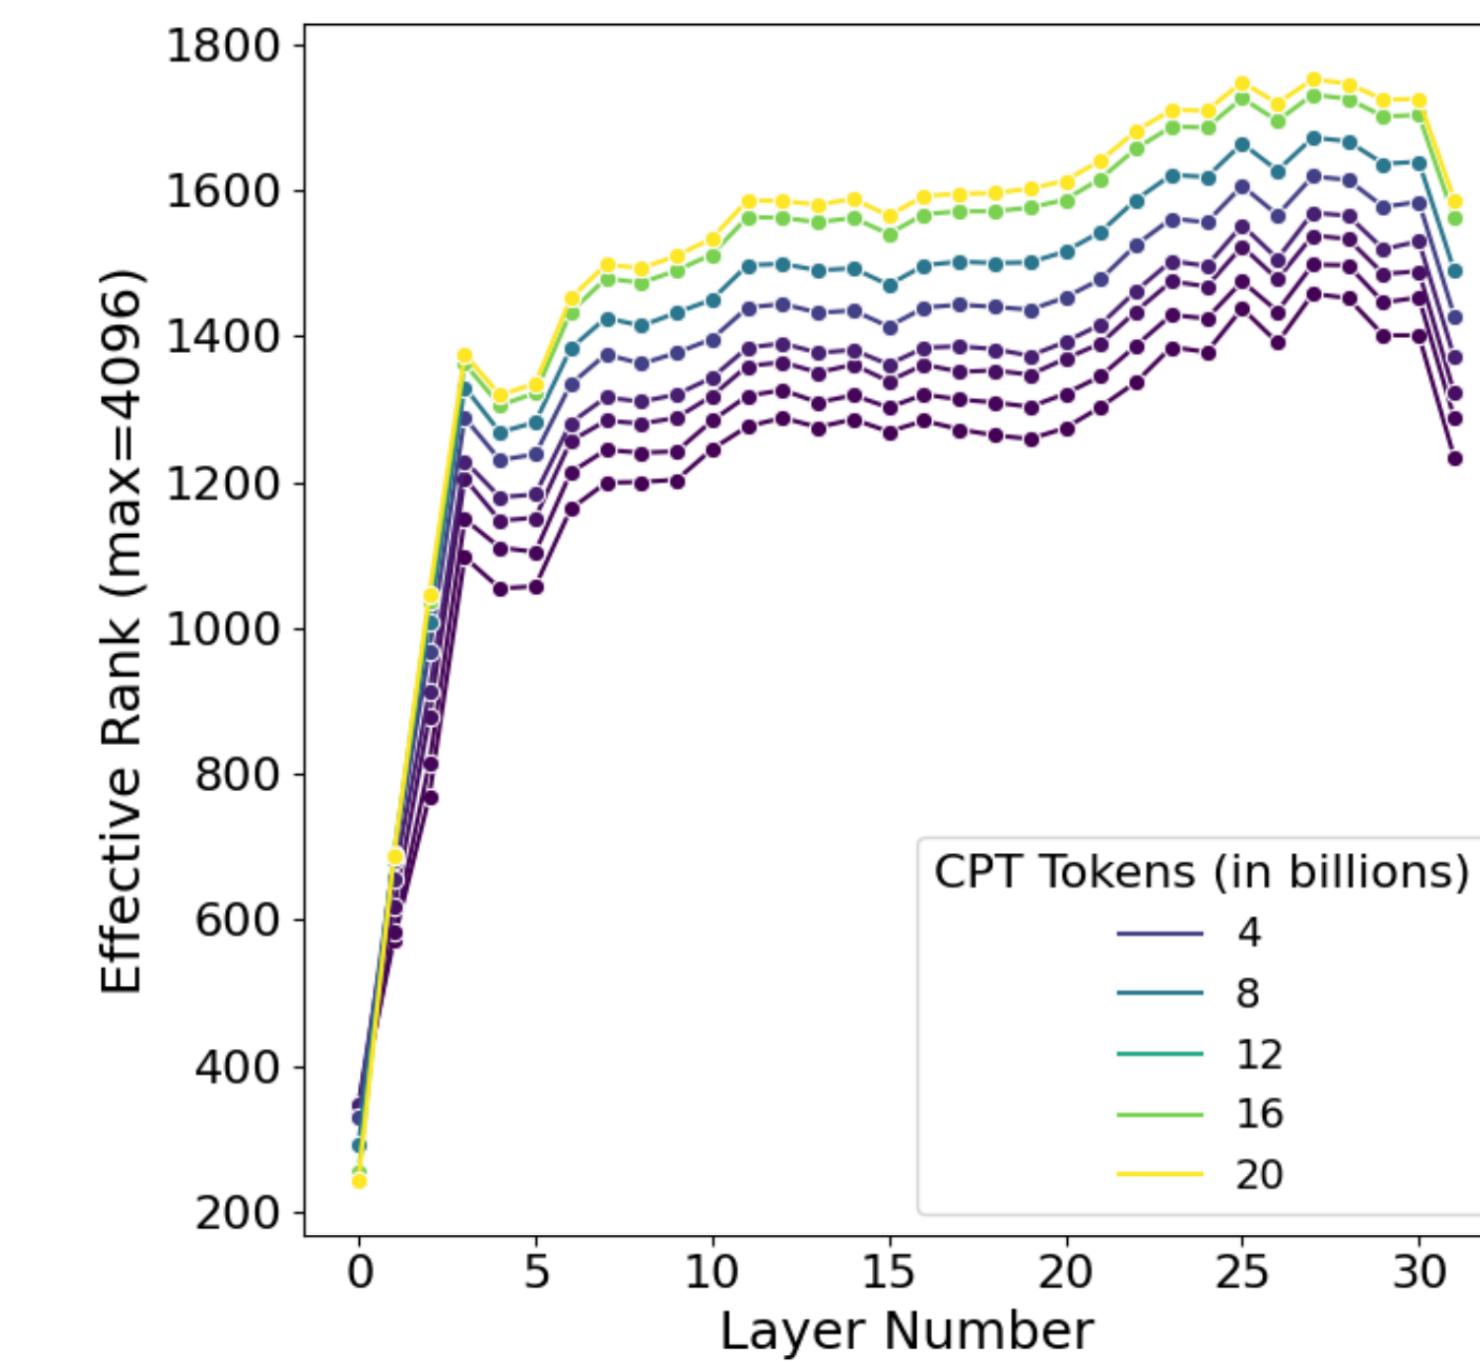
$$r \ll \min(b, a)$$

Low-Rank Adaptation (LoRA)

Low-rank constraint limits the expressive capacity of LoRA



LoRA still lags behind full fine-tuning (FFT)

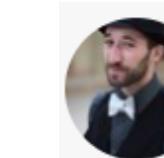


Effective rank of FFT is large

Low-Rank Adaptation (LoRA)

LoRA can be further compressed

- Many related works demonstrate that LoRA can be further compressed by
 - [2] 7.76%
 - [3] 12.5%
 - [4] 50%
 - [5] 3%
- without performance loss.



Leshem Choshen 🤖😊 @LChoshen · Jul 9

...
LoRAs have a lot in similar.
So one can compress (+-SVD with unique s) them together, serve efficiently or understand their shared spaces



Rickard Brüel Gabrielsson @RickardGabriels · Jul 9

Replies to @RickardGabriels
Our work enhances the serving of large language models (LLMs) by efficiently compressing multiple low-rank adapters (LoRAs). We developed linear algebra methods to compress these adapters without significant loss in performance, achieving major throughput ...
[Show more](#)

Compressing 100 LoRAs by sharing their subspaces won't compromise performance.

[2] Kopiczko, Dawid J., Tijmen Blankevoort, and Yuki M. Asano. "Vera: Vector-based random matrix adaptation." ICLR 2023

[3] Renduchintala, Adithya, Tugrul Konuk, and Oleksii Kuchaiev. "Tied-LoRA: Enhancing parameter efficiency of LoRA with weight tying." NAACL 2024

[4] Song, Yurun, et al. "ShareLoRA: Parameter Efficient and Robust Large Language Model Fine-tuning via Shared Low-Rank Adaptation." arXiv 2024

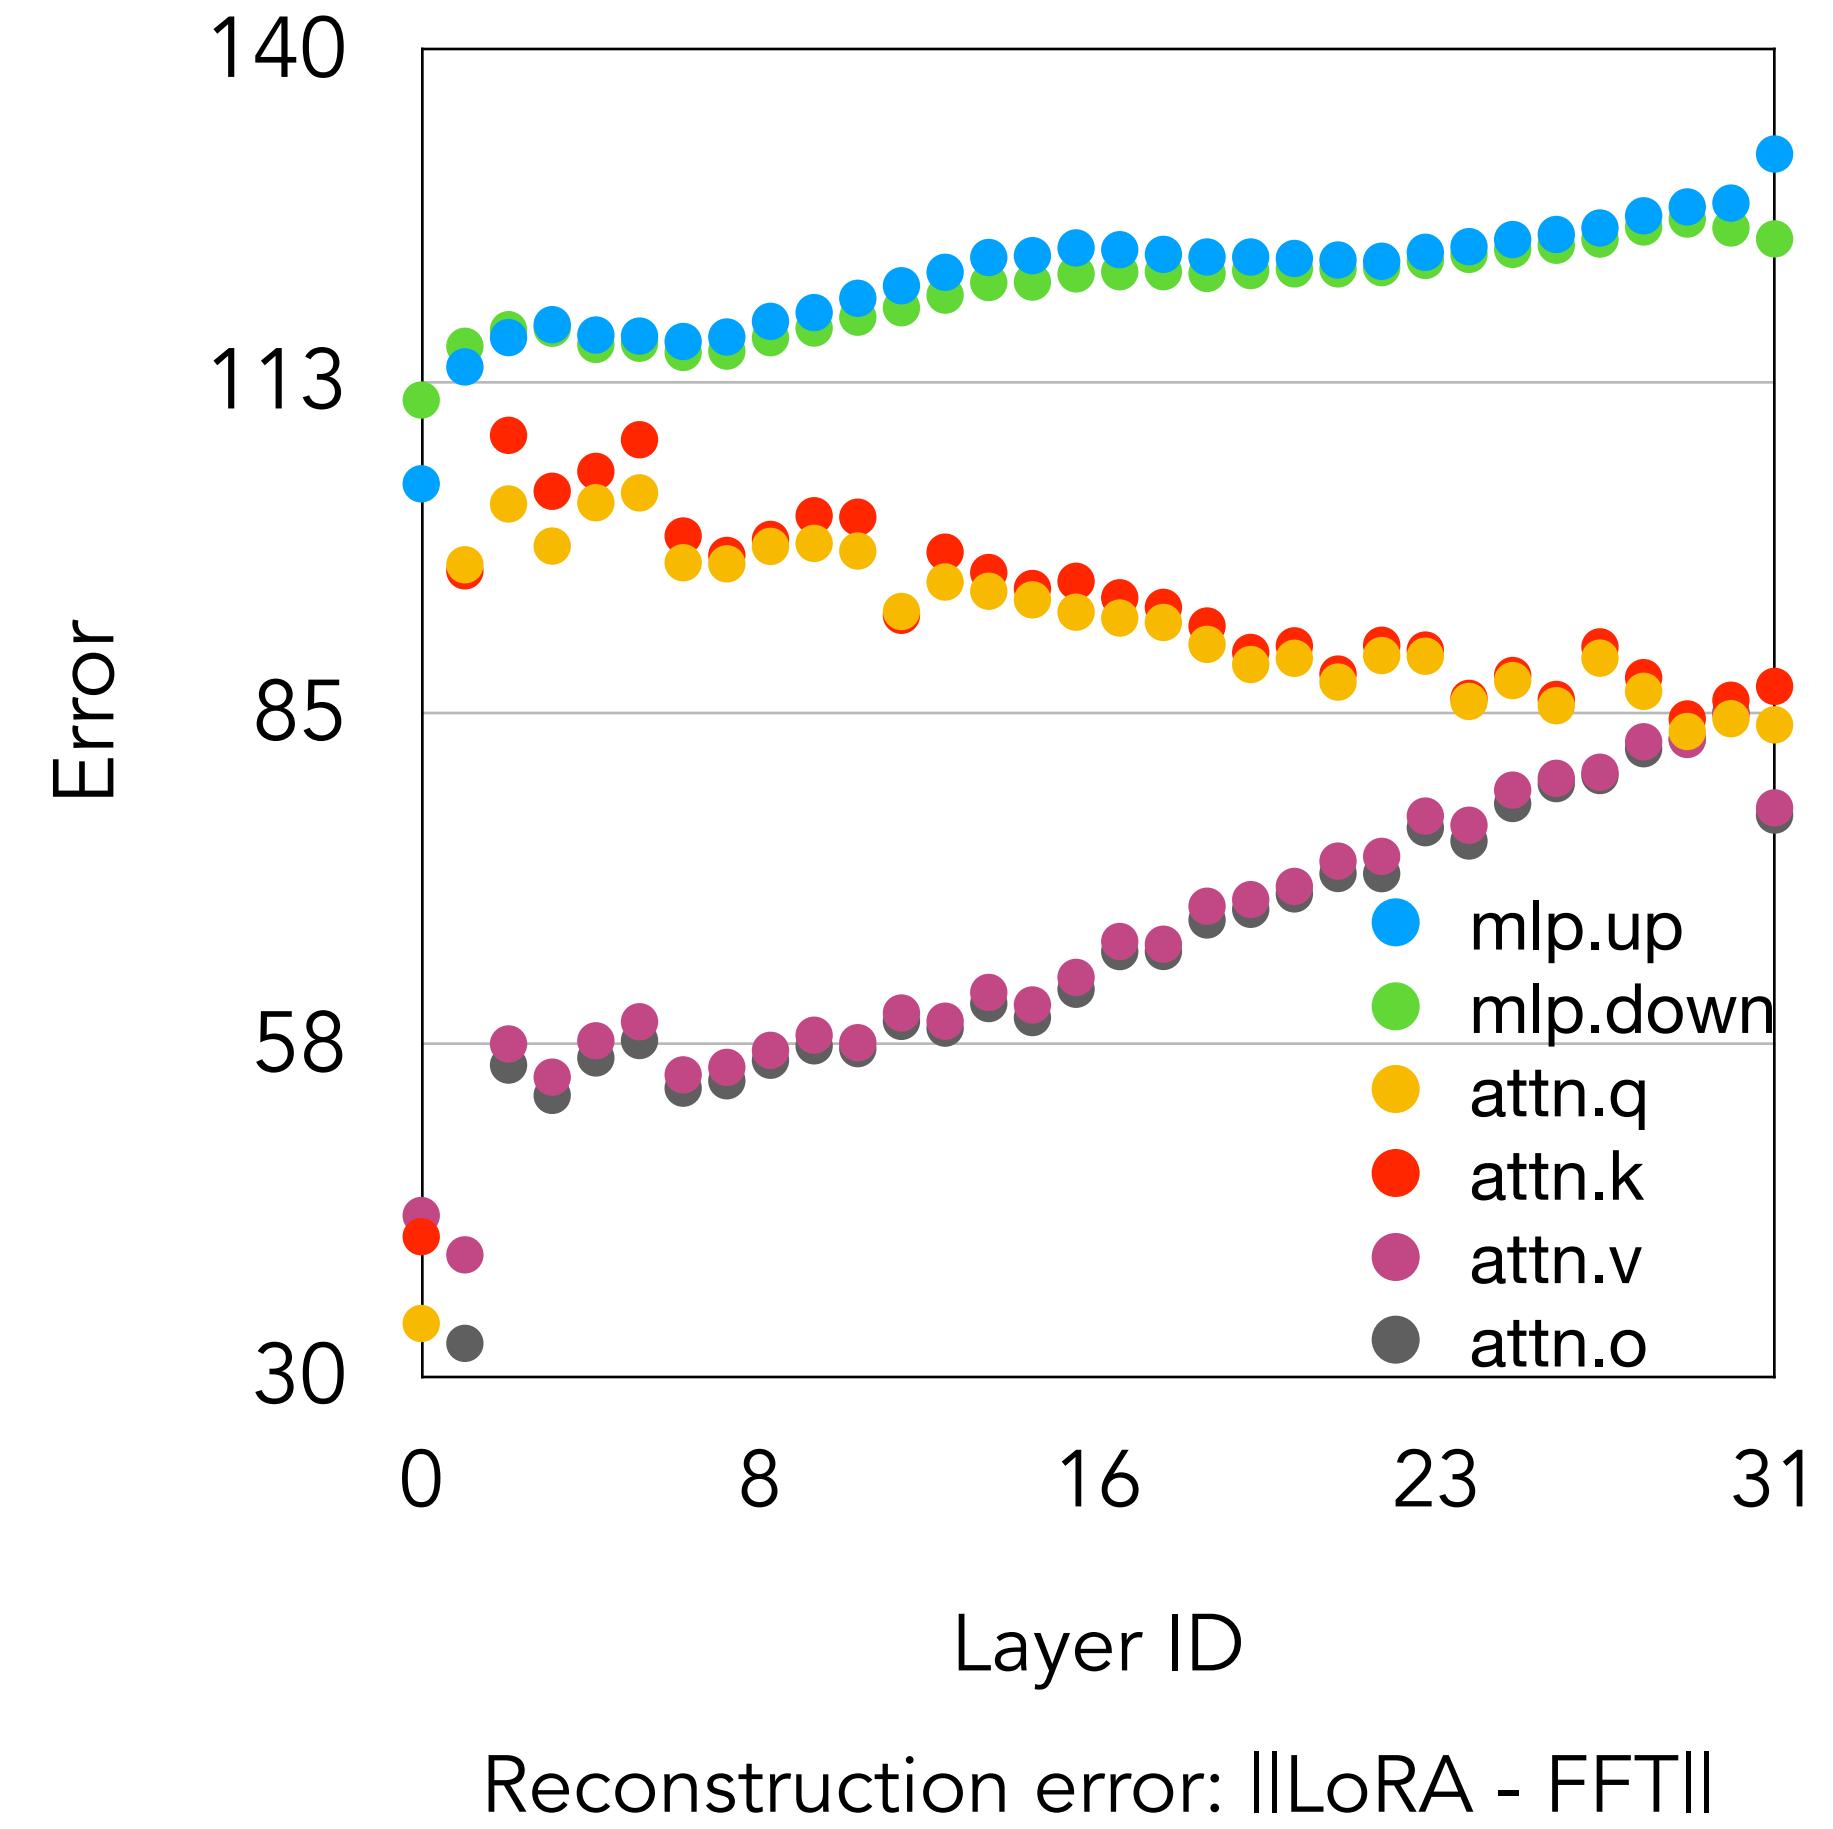
[5] Li, Yang, Shaobo Han, and Shihao Ji. "VB-LoRA: extreme parameter efficient fine-tuning with vector banks." NeurIPS 2024

LoRA is underutilized.

Low-Rank Adaptation (LoRA)

Why LoRA is under-utilized?

- Different components and layers require different levels of expressive capability.
- LoRA adopts an average allocation strategy.



Rank-Sharing Low-Rank Adaptation (RaSA)

Partial rank-sharing across layers

- Split the matrices \mathbf{B}_i and \mathbf{A}_i into layer-specific parts $(\tilde{\mathbf{B}}_i, \tilde{\mathbf{A}}_i)$ and layer-shared parts $(\hat{\mathbf{B}}_i, \hat{\mathbf{A}}_i)$

$$\mathbf{B}_i = \begin{bmatrix} \underbrace{\tilde{\mathbf{B}}_i}_{\mathbb{R}^{b \times (r-k)}} & \underbrace{\hat{\mathbf{B}}_i}_{\mathbb{R}^{b \times k}} \end{bmatrix}, \quad \mathbf{A}_i = \begin{bmatrix} \underbrace{\tilde{\mathbf{A}}_i^T}_{\mathbb{R}^{a \times (r-k)}} & \underbrace{\hat{\mathbf{A}}_i^T}_{\mathbb{R}^{a \times k}} \end{bmatrix}^T$$

- Concatenate all layer-shared parts across layers to form shared rank pools $(\mathbf{B}_S$ and $\mathbf{A}_S)$

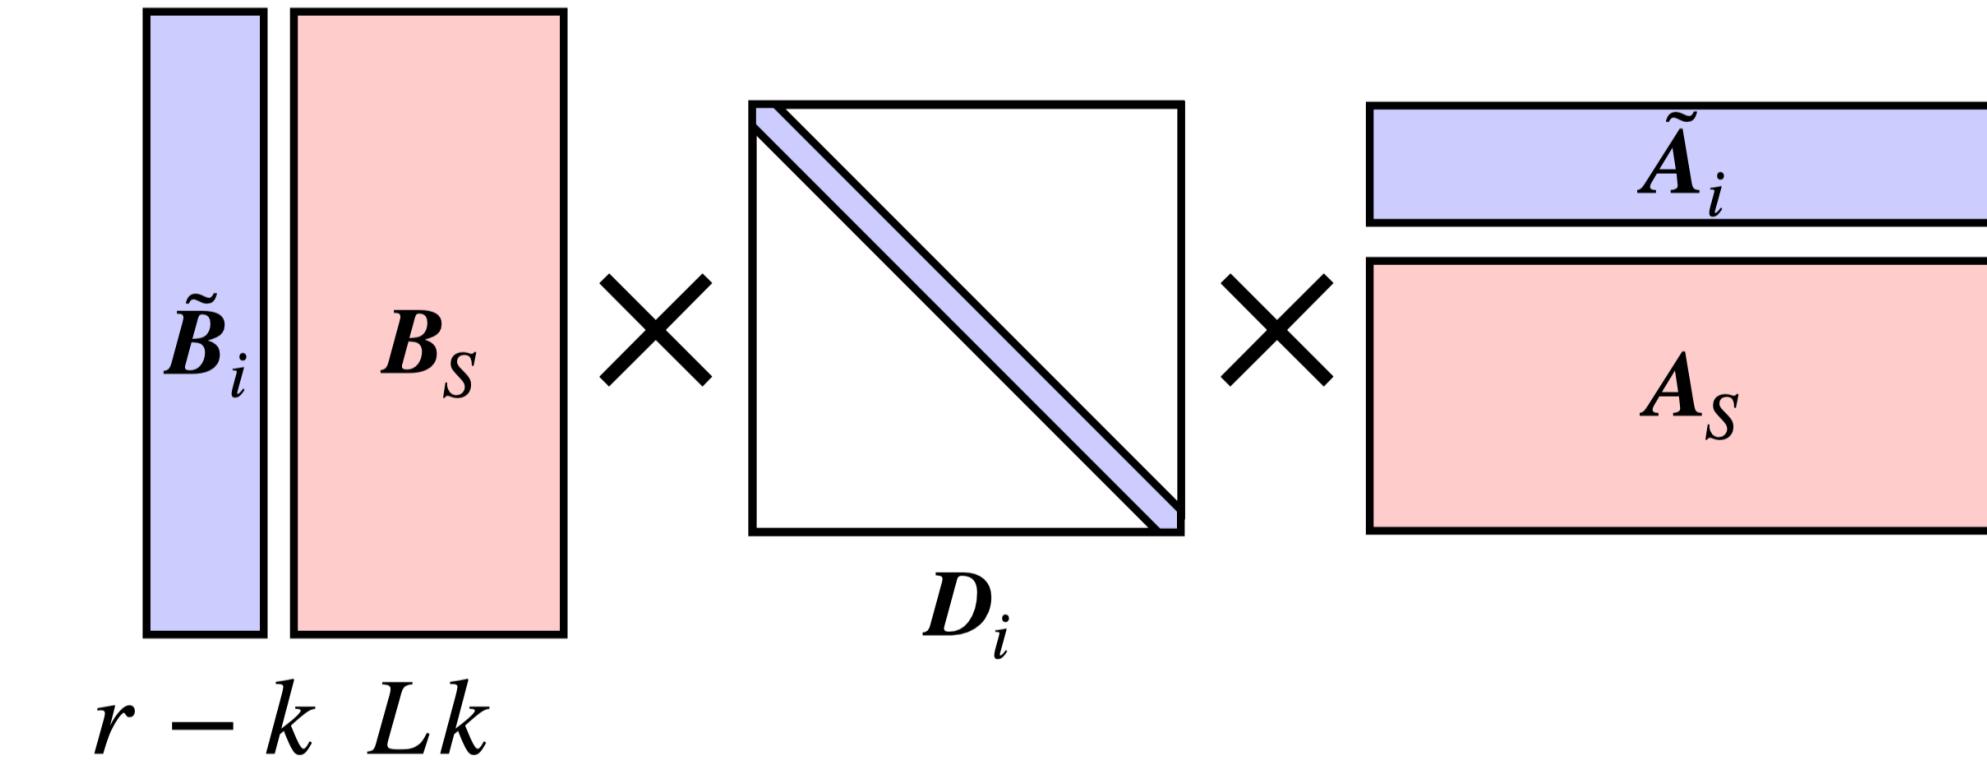
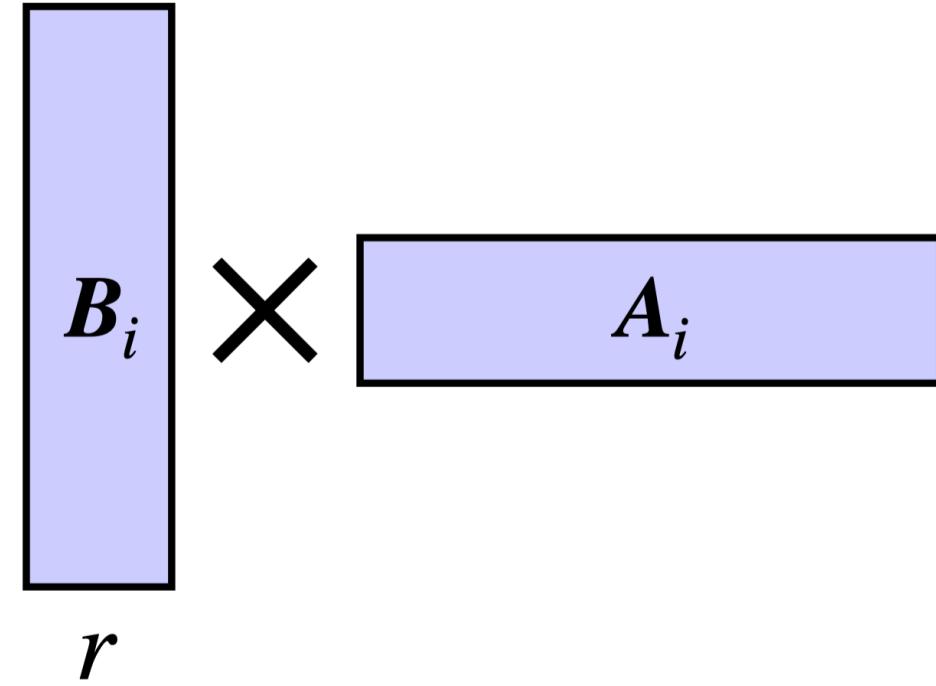
$$\mathbf{B}_S = [\hat{\mathbf{B}}_1 \quad \hat{\mathbf{B}}_2 \quad \dots \quad \hat{\mathbf{B}}_L] \in \mathbb{R}^{b \times (L \times k)}, \quad \mathbf{A}_S = [\hat{\mathbf{A}}_1^T \quad \hat{\mathbf{A}}_2^T \quad \dots \quad \hat{\mathbf{A}}_L^T]^T \in \mathbb{R}^{(L \times k) \times a}$$

- Update of layer-i

$$\begin{aligned} \mathbf{W}_i + \Delta \mathbf{W}_i &= \mathbf{W}_i + \frac{\alpha}{r} (\tilde{\mathbf{B}}_i \tilde{\mathbf{A}}_i + \mathbf{B}_S \mathbf{A}_S) \\ &= \mathbf{W}_i + [\tilde{\mathbf{B}}_i \quad \mathbf{B}_S] \operatorname{diag}\left(\frac{\alpha}{r}\right) \begin{bmatrix} \tilde{\mathbf{A}}_i \\ \mathbf{A}_S \end{bmatrix} \end{aligned}$$

Rank-Sharing Low-Rank Adaptation (RaSA)

Comparison between LoRA and RaSA



$$W_i + \Delta W_i = W_i + \frac{\alpha}{r} B_i A_i \quad (B_i \in \mathbb{R}^{b \times r}, A_i \in \mathbb{R}^{r \times a})$$

LoRA

$$W_i + \Delta W_i = W_i + \underbrace{[\tilde{B}_i \quad B_S]}_{\mathbb{R}^{b \times (r-k+Lk)}} D_i \underbrace{\begin{bmatrix} \tilde{A}_i \\ A_S \end{bmatrix}}_{\mathbb{R}^{(r-k+Lk) \times a}}$$

RaSA

- $r \Rightarrow r-k+Lk$
- extra parameters (0.01%)

Reconstruction Error Analysis

Minimum Reconstruction Error

- We compare their abilities to reconstruct a set of high-rank matrices $\{\mathbf{M}_i\}_{i \in [L]}$, $\text{rank}(\mathbf{M}_i) = R > r$

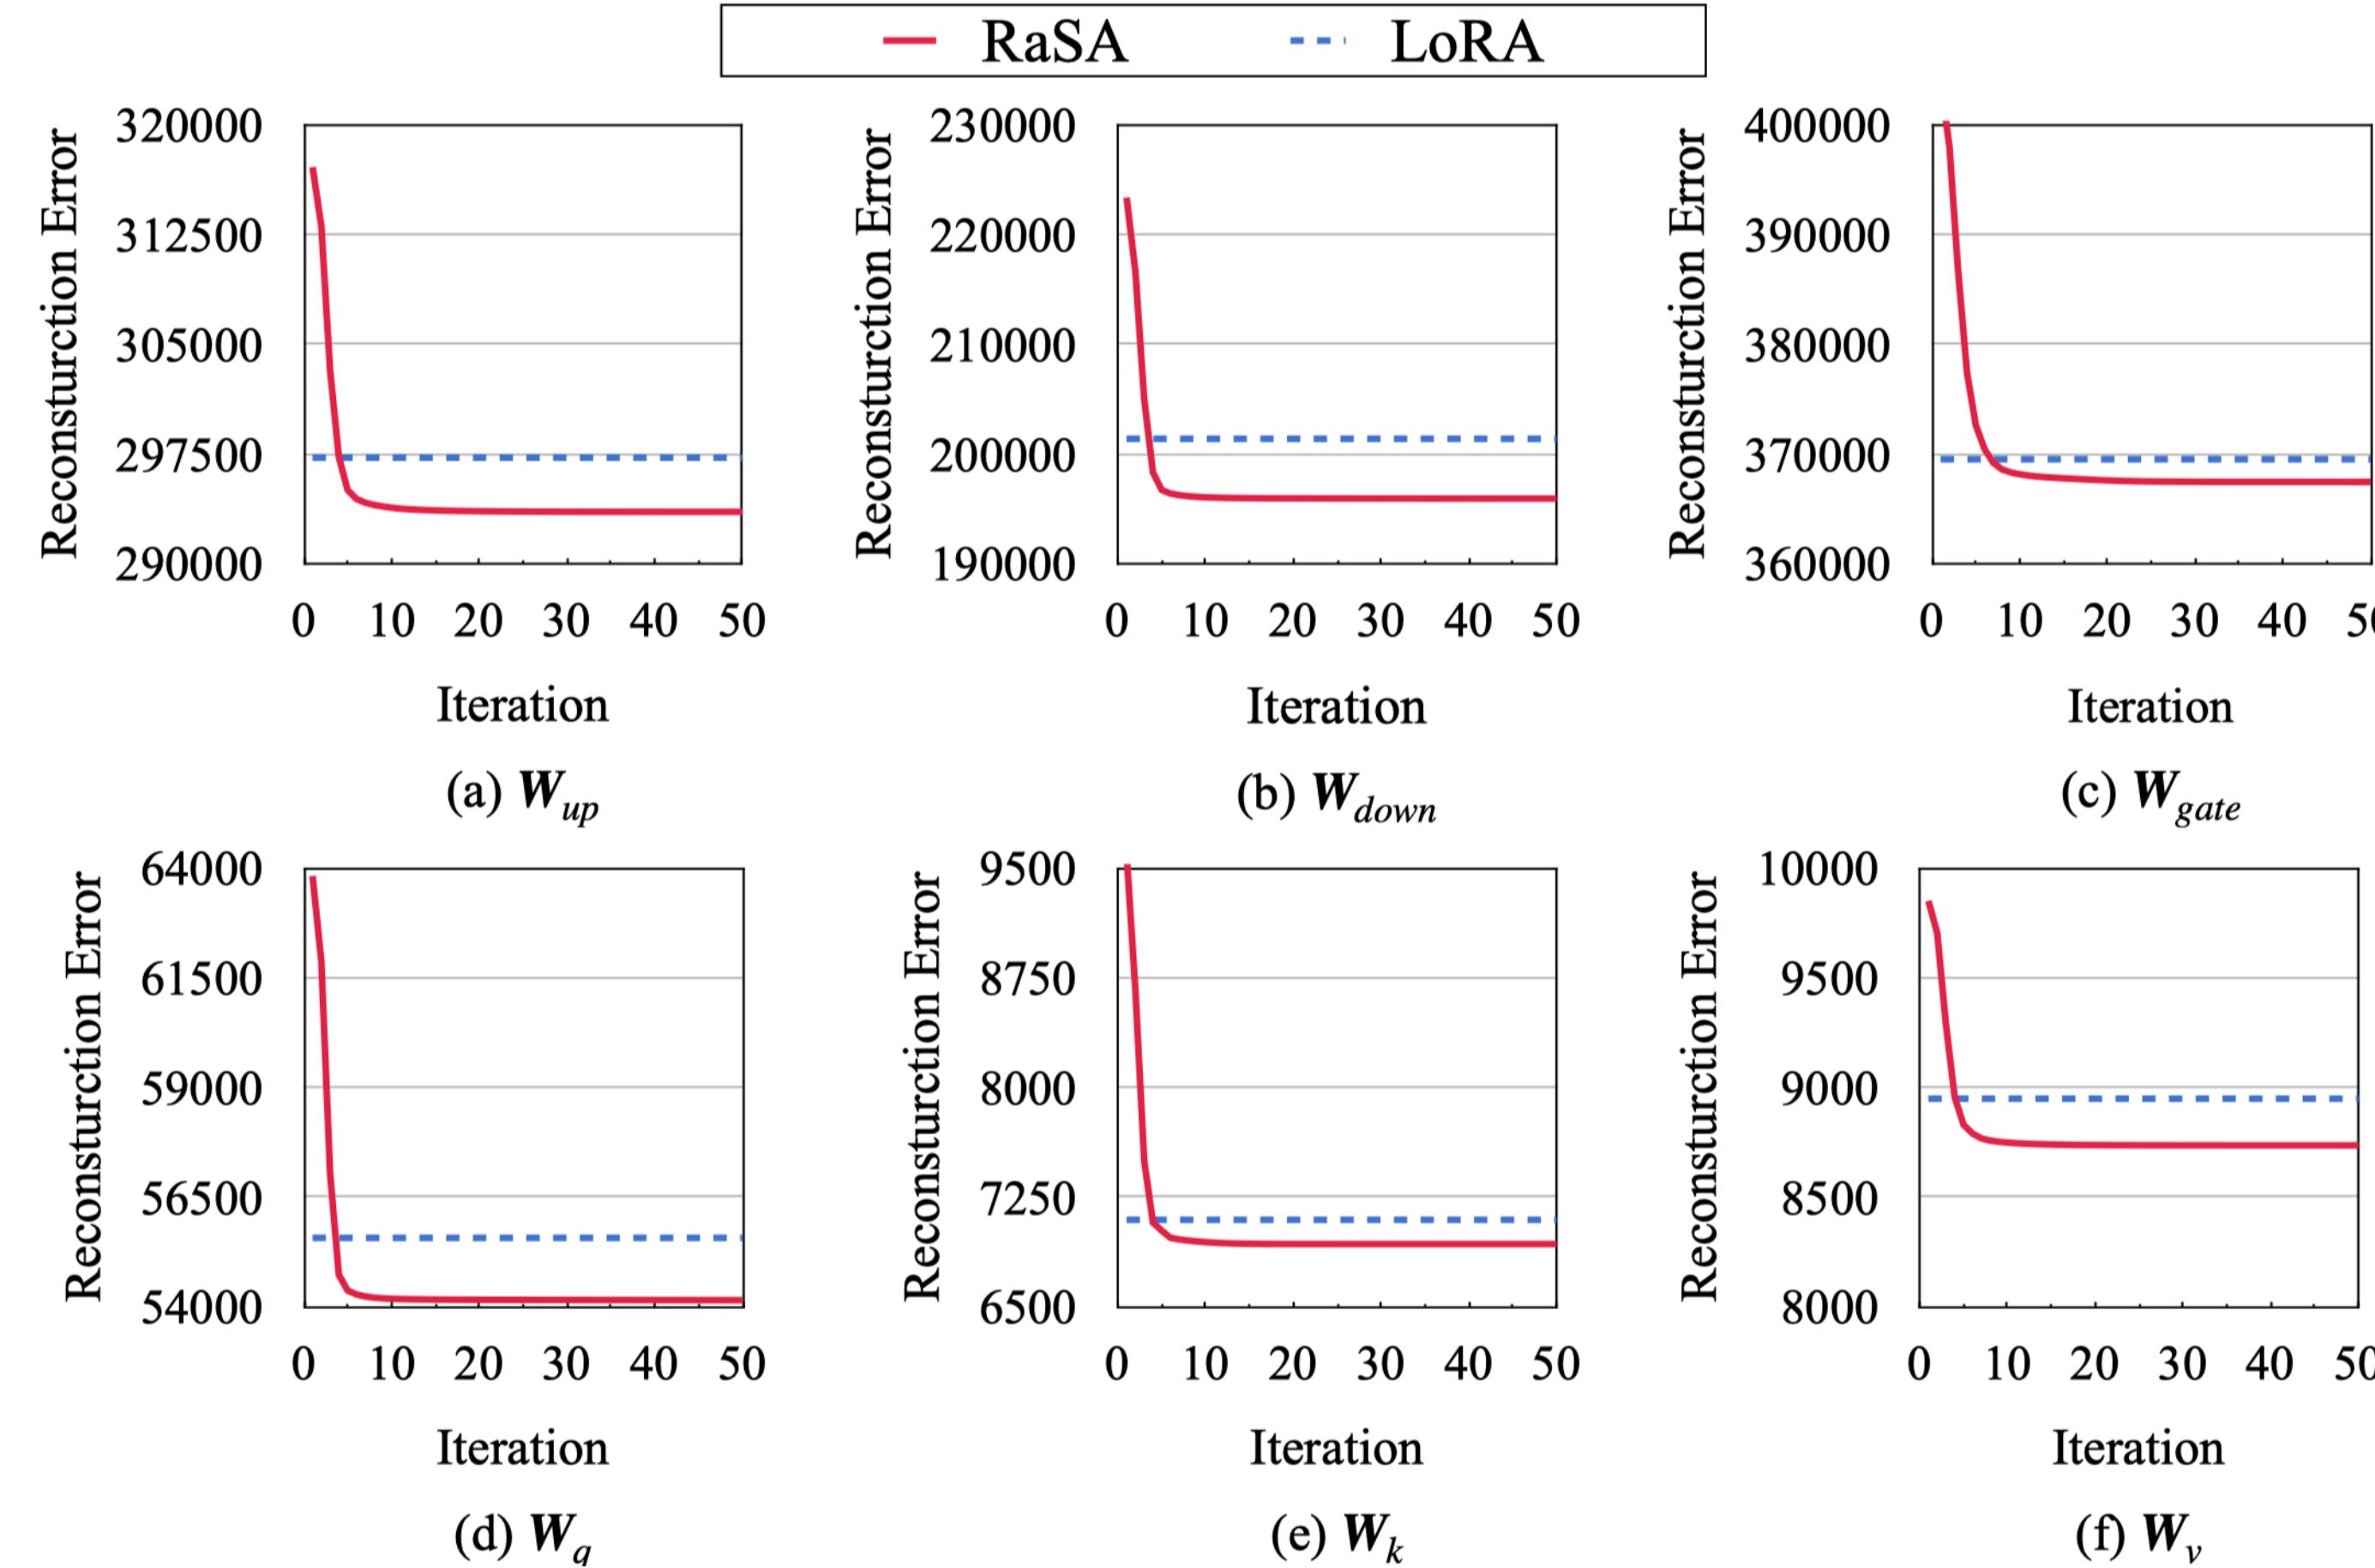
$$e_{\text{lora}} = \min_{\mathbf{B}_i, \mathbf{A}_i} \sum_{i=1}^L \|\mathbf{M}_i - \mathbf{B}_i \mathbf{A}_i\|_F^2$$

$$e_{\text{rasa}(k)} = \min_{\tilde{\mathbf{B}}_i, \tilde{\mathbf{A}}_i, \mathbf{B}_S, \mathbf{A}_S, \mathbf{D}_i} \sum_{i=1}^L \|\mathbf{M}_i - (\tilde{\mathbf{B}}_i \tilde{\mathbf{A}}_i + \mathbf{B}_S \mathbf{D}_i \mathbf{A}_S)\|_F^2$$

- We prove $e_{\text{rasa}(k)} \leq e_{\text{lora}}$ (Theorem 3.1)

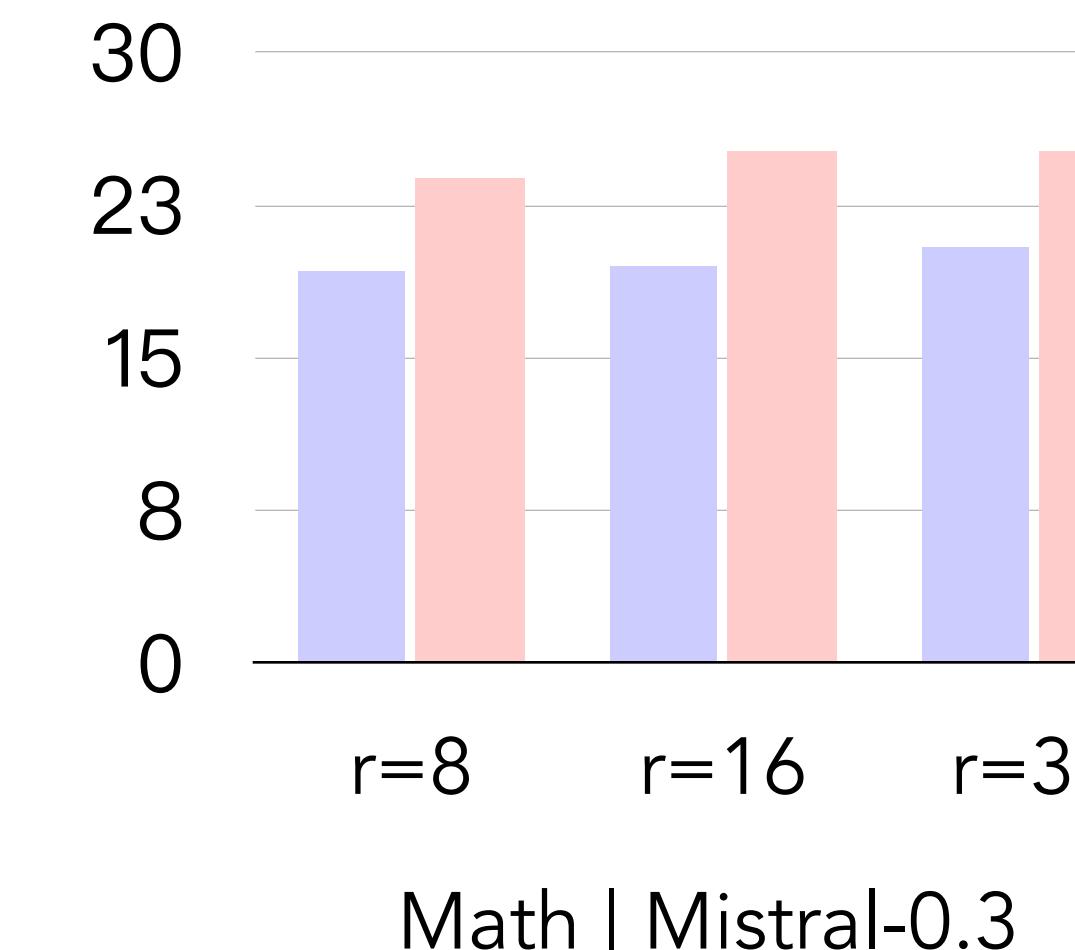
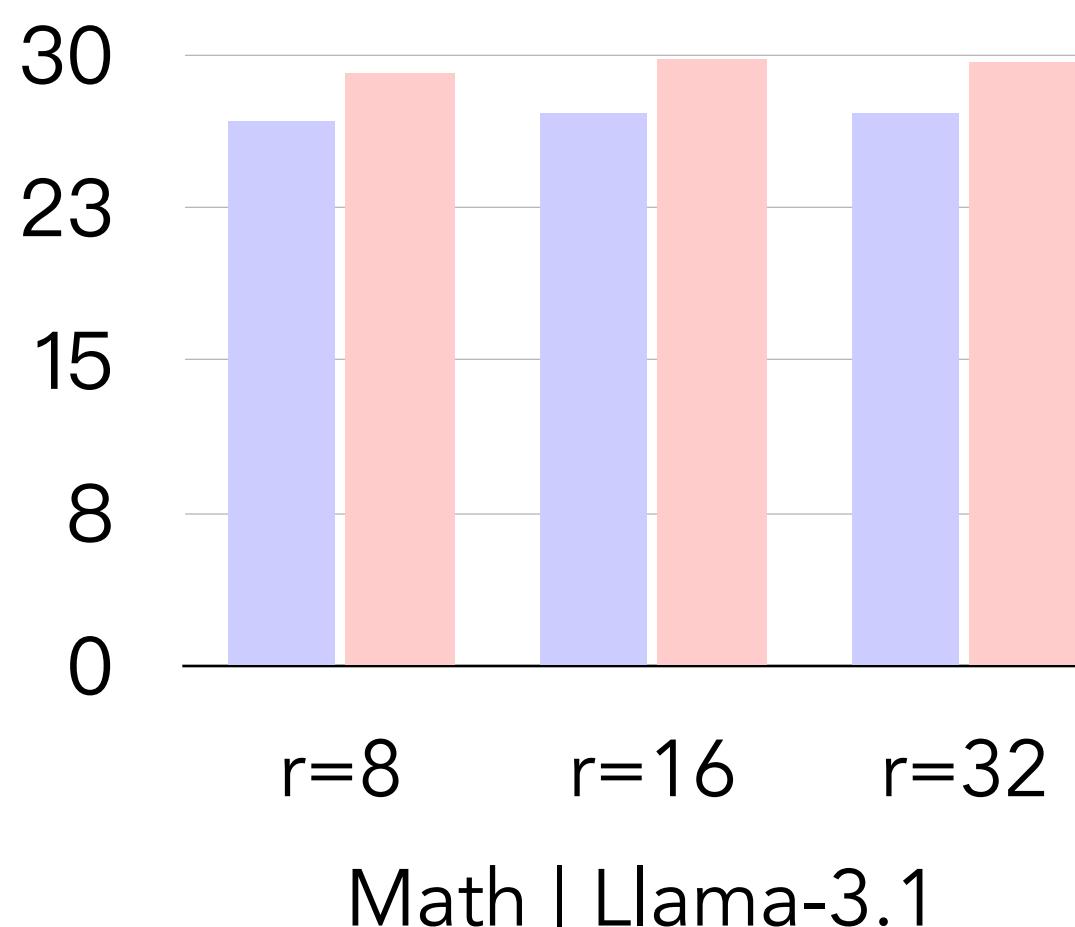
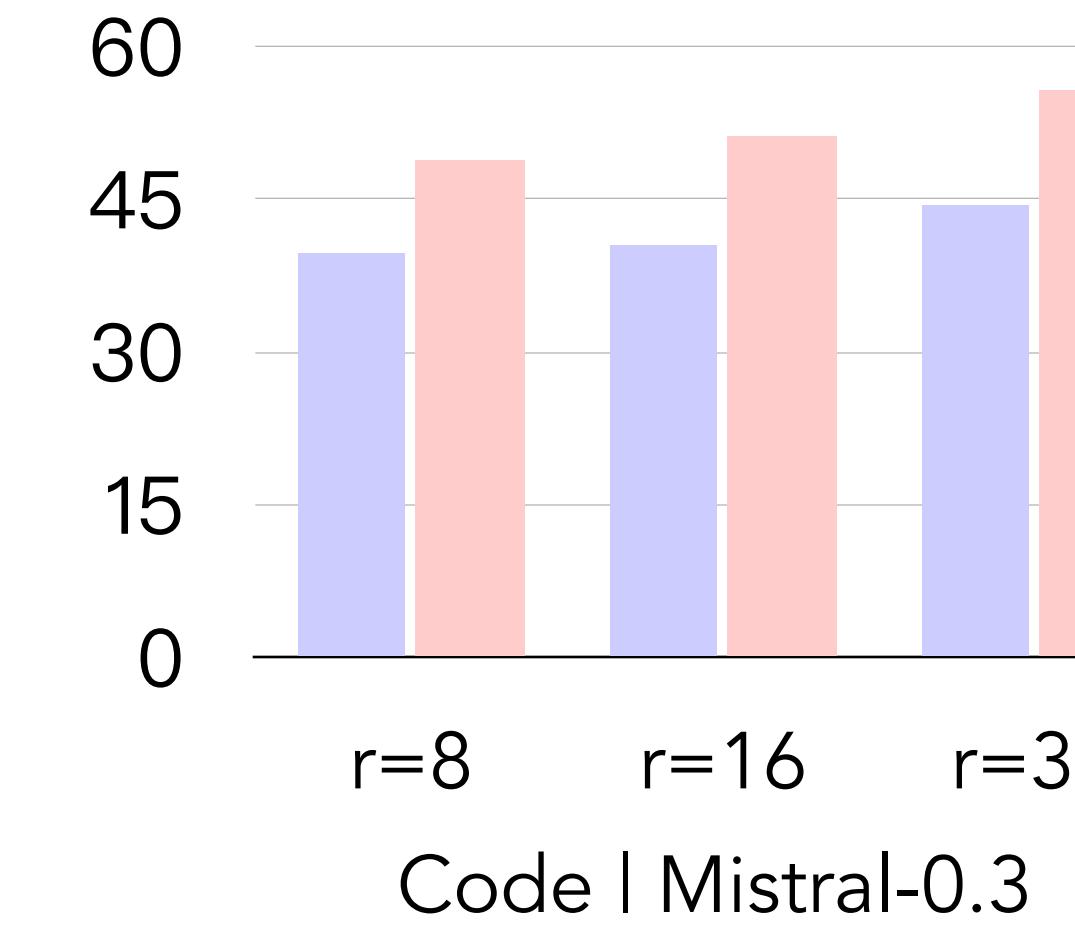
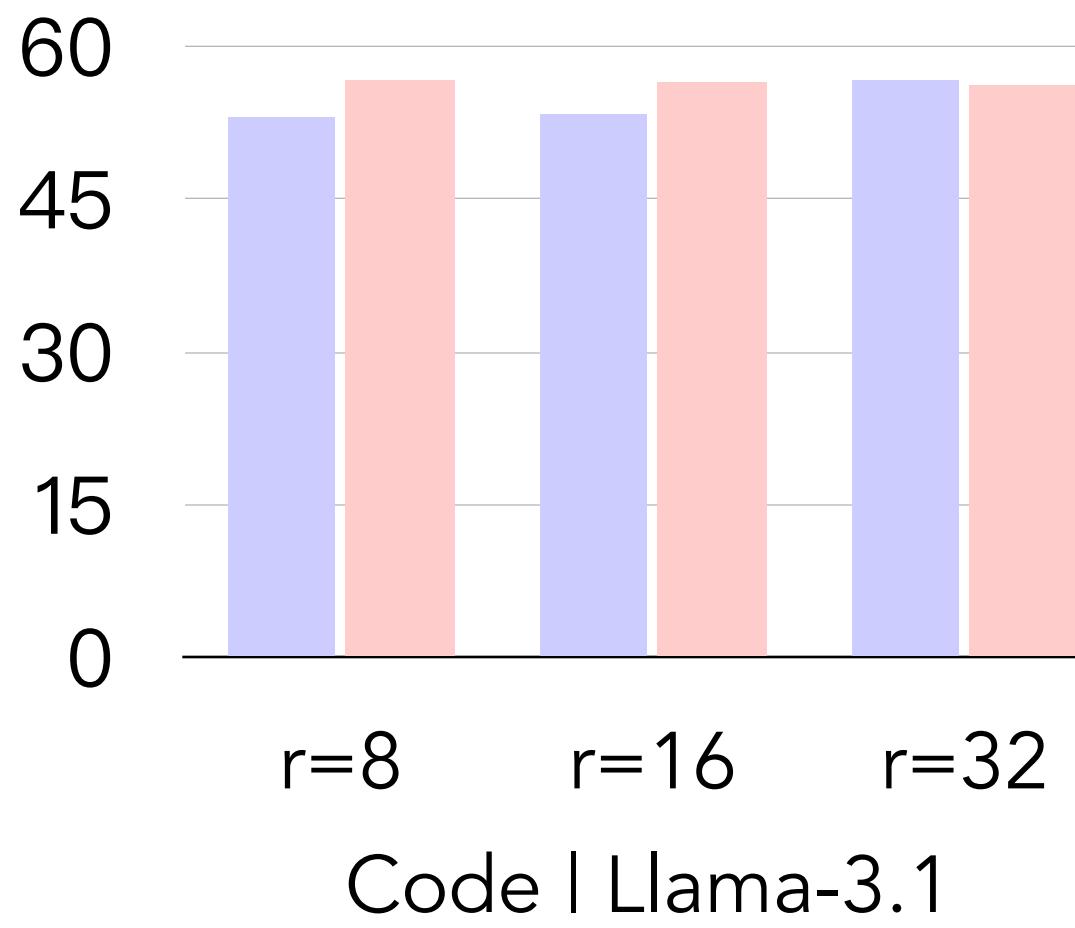
Reconstruction Error Analysis

Empirical Analysis

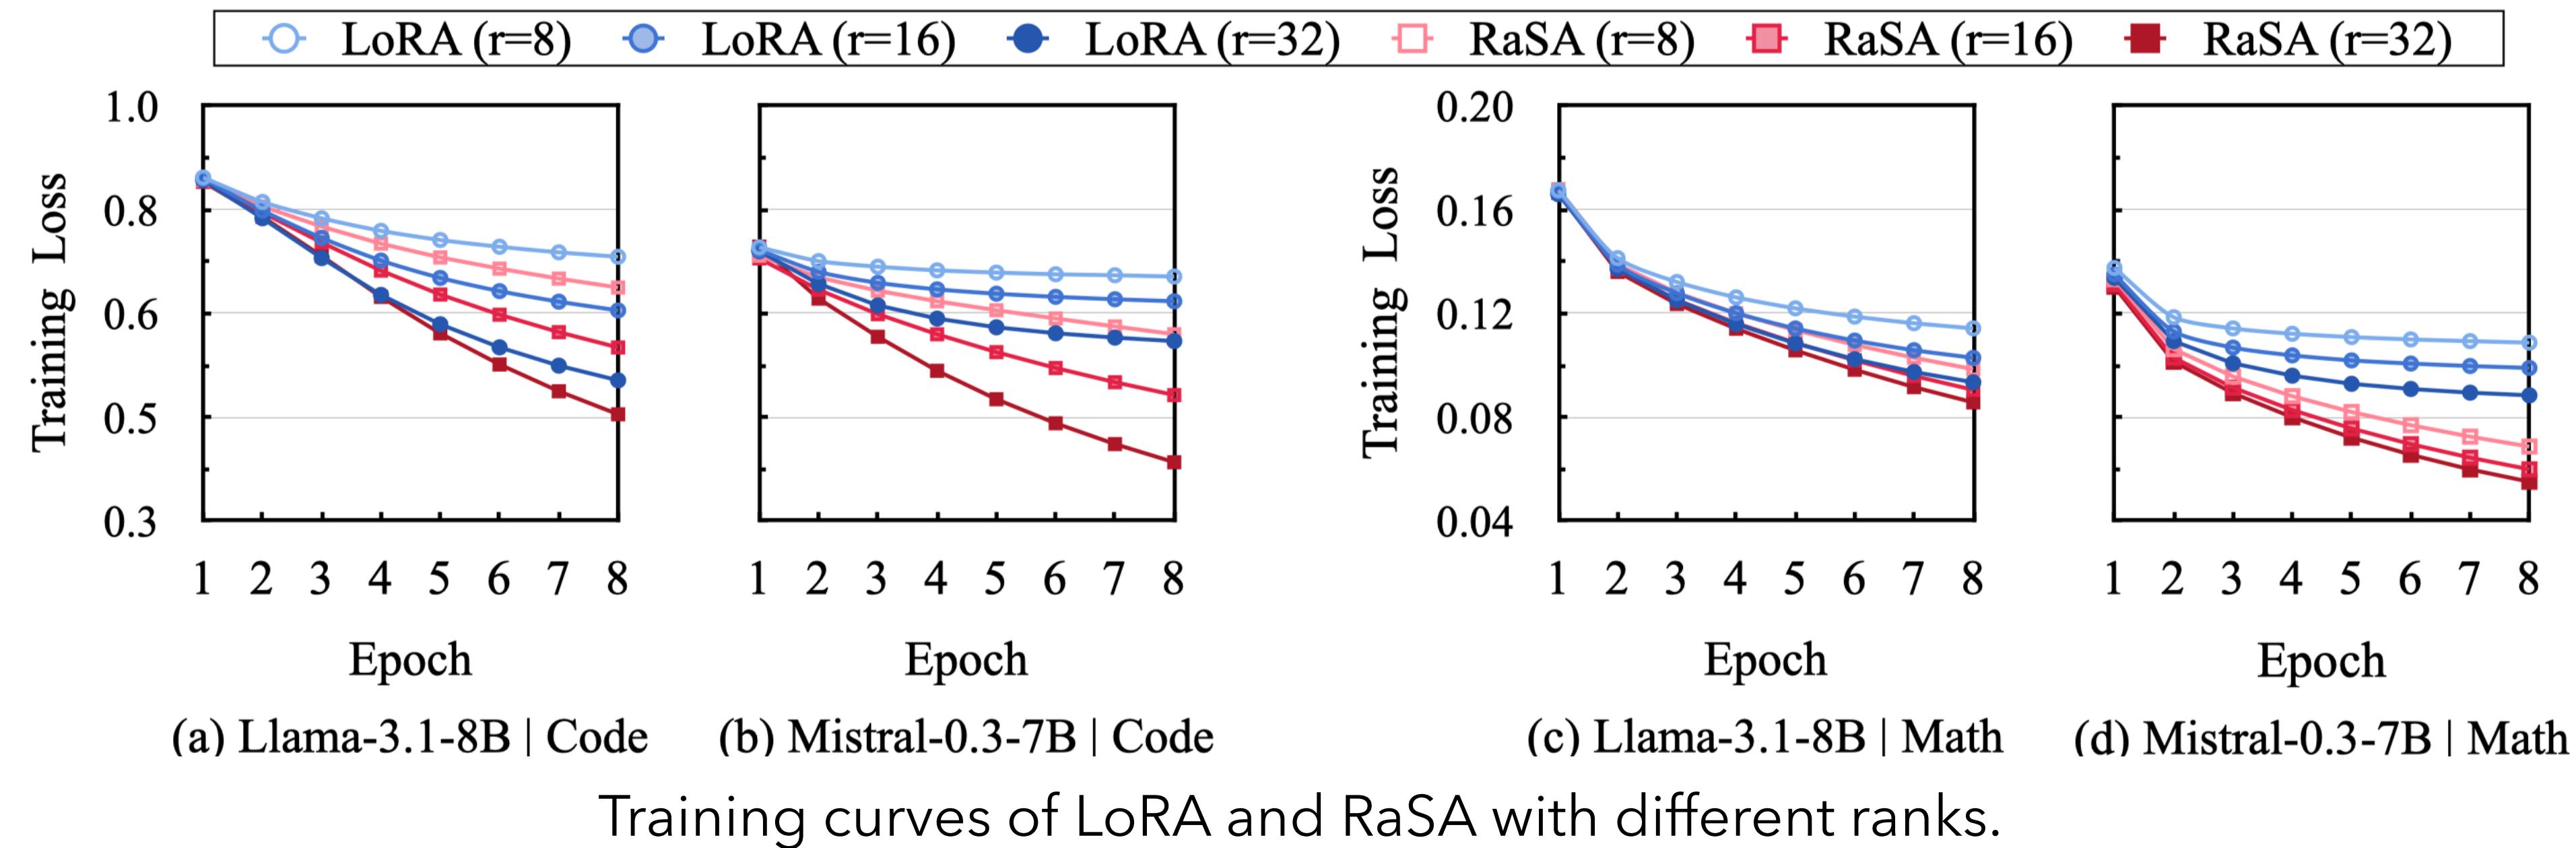


RaSA requires ~10 iterations to achieve a significantly lower reconstruction error than LoRA's minimum.

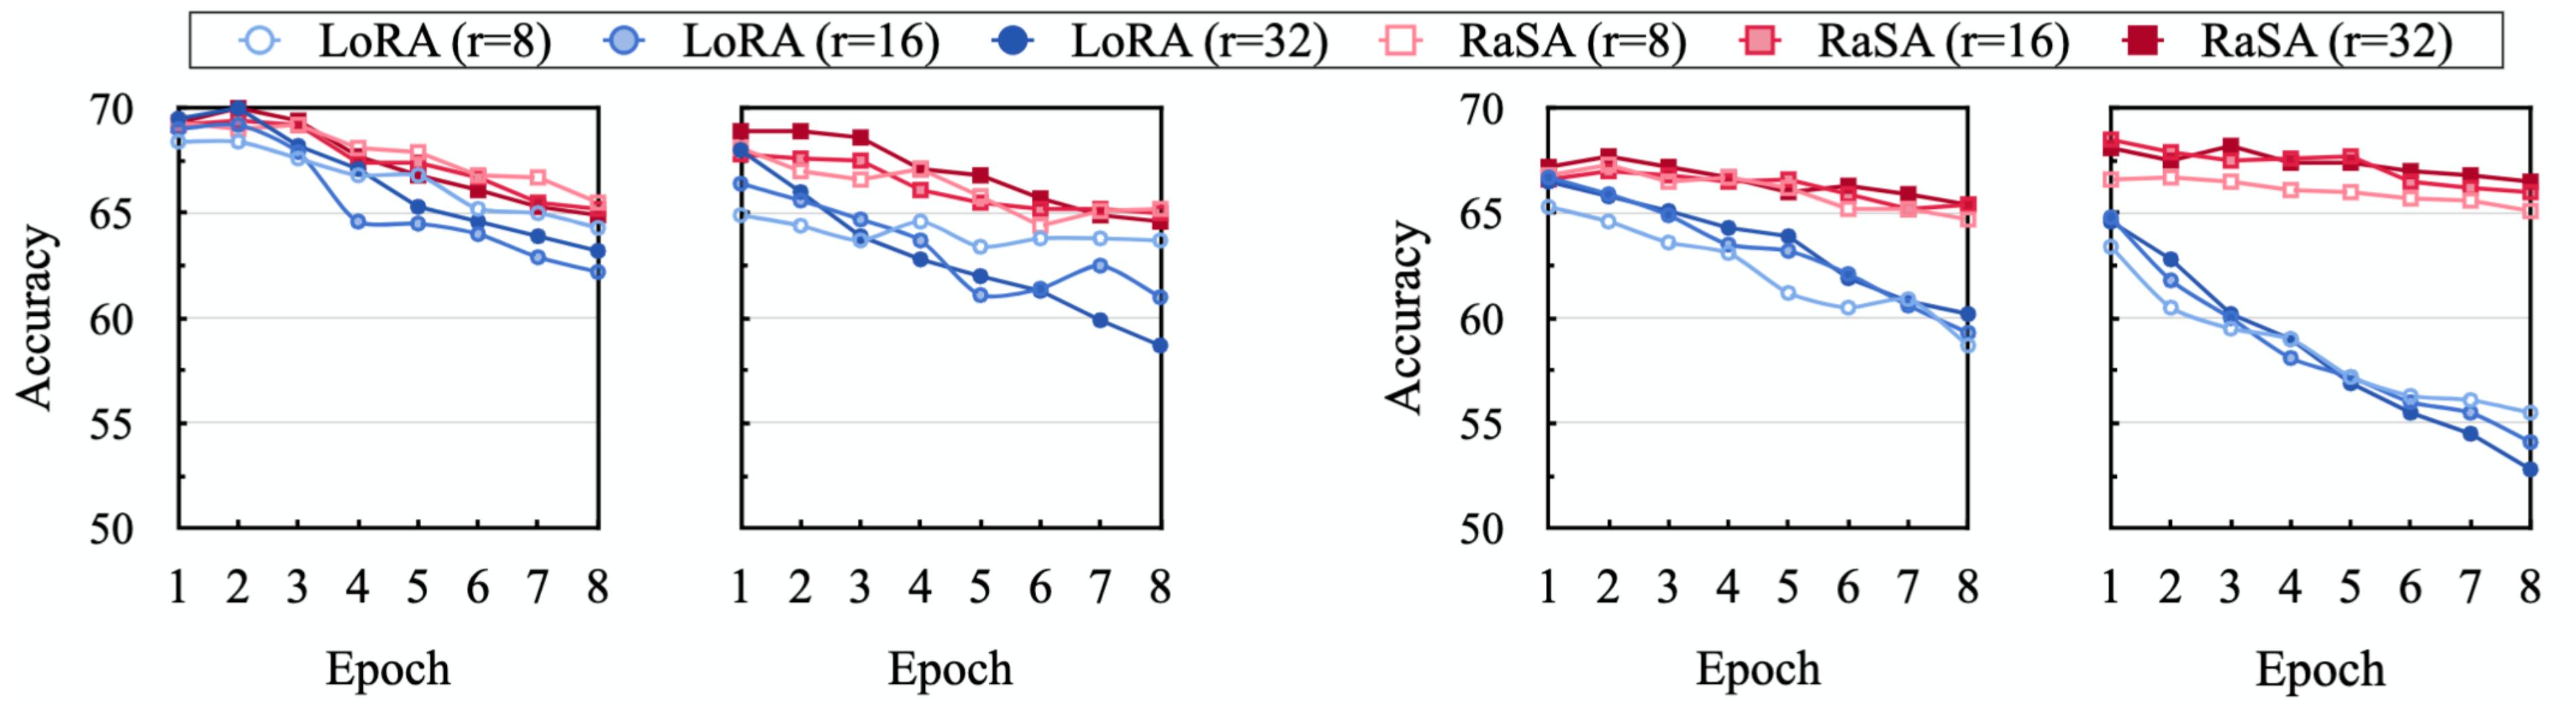
Main Results



RaSA learns more and faster than LoRA



RaSA forgets less than LoRA



(a) Llama-3.1-8B | Code (b) Mistral-0.3-7B | Code

(c) Llama-3.1-8B | Math (d) Mistral-0.3-7B | Math

Y-axis shows the average of prediction accuracy on three benchmarks to evaluate model's forgetting. Higher prediction accuracy denotes less forgetting.

Summary

- We propose RaSA, an extension of LoRA by allowing partial rank sharing across layers, which significantly improves the efficiency and expressiveness.
- We provide a comprehensive analysis – both theoretical and empirical – showcasing RaSA's superior capacity for matrix reconstruction and its resultant improved performance on downstream tasks.

Thank You